

IESTI01 – TinyML

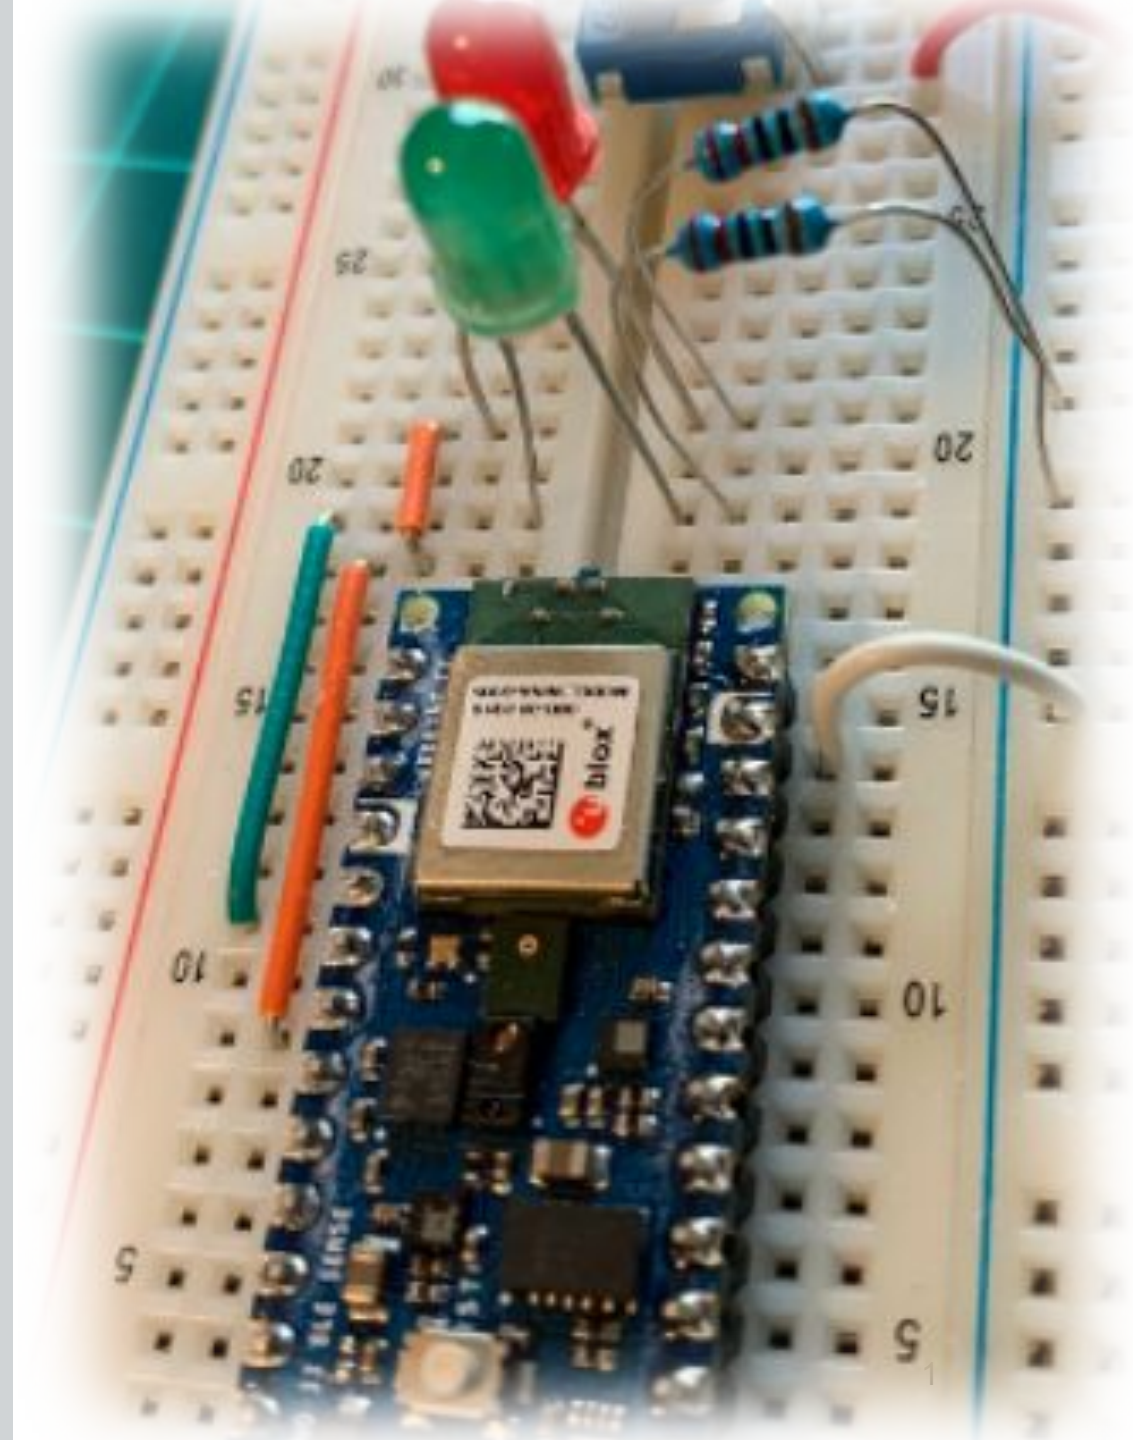
Embedded Machine Learning

14. Fundamentals wrap-up and Application's preview



Prof. Marcelo Rovai

UNIFEI



Tiny Machine Learning (TinyML)


What we learned so far

What is Tiny Machine Learning (**TinyML**)?

- Fast-growing field of **machine learning**
- Algorithms, **hardware, and software**
- **On-device** sensor data analytics
- Extreme **low power** consumption
- **Always-on ML** use-cases
- **Battery**-operated devices


What we already learned?

Part 1



So far in the Part 1, we introduced ML with TensorFlow. Was all about talking about what is the **language of machine learning**.

Total Recall from **Part 1**



“Language” for Part 1

Neural Network

Gradient Descent

Loss Function

Training Data

Training

Validation Data

Inference

Test Data

Features

Classification

Filters

Overfitting

Regression

Kernels

Data augmentation

Responsible AI

CNNs

DNNs

Preprocessing

“Language” for Part 1

Training Data

Neural Network

Training

Validation Data

Gradient Descent

Inference

Test Data

Loss Function

Features

Classification

Filters

Overfitting

Regression

Kernels

Data augmentation

Responsible AI

CNNs

DNNs

Preprocessing

“Language” for Part 1

Training Data

Neural Network

Training

Validation Data

Gradient Descent

Inference

Test Data

Loss Function

Features

Classification

Filters

Overfitting

Kernels

Regression

Data augmentation

CNNs

DNNs

Responsible AI

Preprocessing

“Language” for Part 1

Training Data

Neural Network

Training

Validation Data

Gradient Descent

Inference

Test Data

Loss Function

Features

Classification

Filters

Overfitting

Kernels

Regression

CNNs

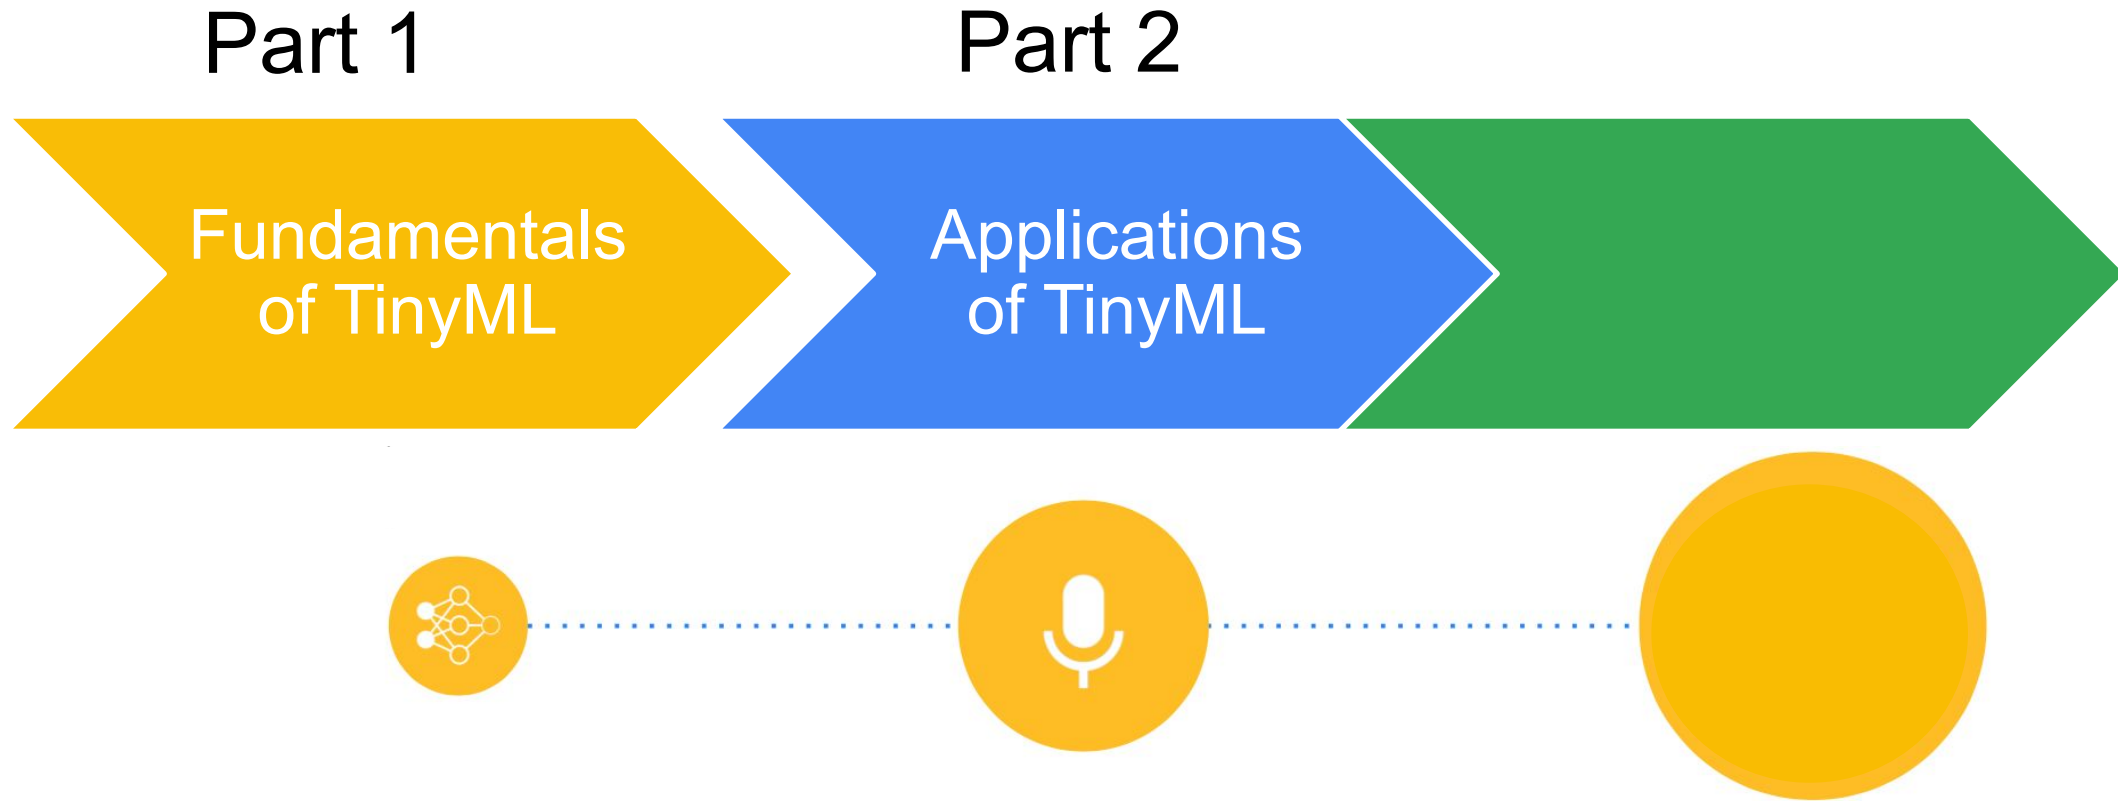
Data augmentation

DNNs

Responsible AI

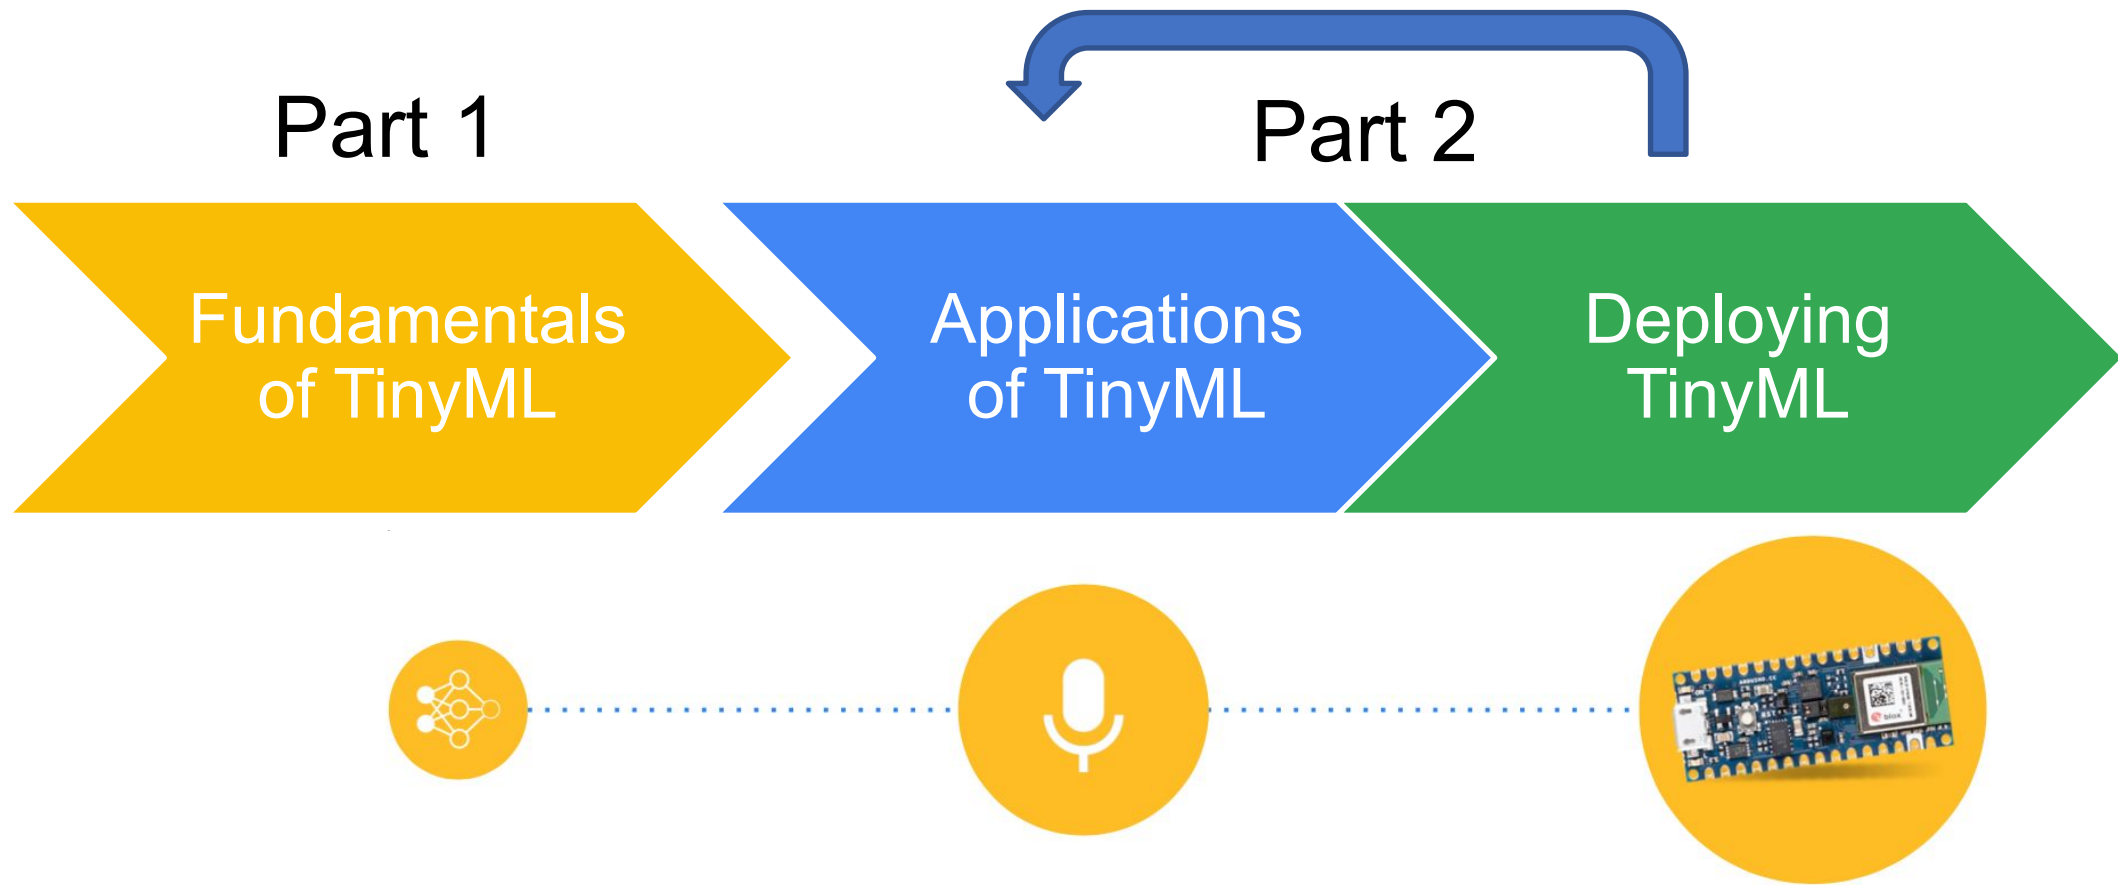
Preprocessing

What we will learn?




In Part 2, we will get a sneak peek into the variety of different TinyML applications, as keyword spotting (“Alexa”), gesture recognition, understand how to leverage the sensors, and so forth.

What we will learn?



In Part 2, we will **also** learn how to deploy models on a real microcontroller. Along the way we will explore the challenges unique to and amplified by TinyML (e.g., preprocessing, post-processing, dealing with resource constraints).



Train a model

Convert
model

Optimize
model

Deploy
model at
Edge

Make
inferences
at Edge







Train a model

Convert
model

Optimize
model

Deploy
model at
Edge

Make
inferences
at Edge



Tiny Machine Learning (TinyML)

Applications

TinyML Application Areas



Home



Office



Industry

TinyML Application Areas



Home



Office



Industry



Questions

- How do we **capture** the data to feed into the neural network?
- How do you **design** the neural network to take in the speech signal?
- What **dataset** does the neural network need to be trained?
- How do we **pre-process** the data for neural network inference?
- How do you **post-process** the neural network output?
- How do you make sure there is no **bias** in the dataset?
- How do you **deploy** this on the microcontroller?

Endpoints Have **Sensors**, Tons of Sensors

Motion Sensors

Gyroscope, Radar,
Accelerometer

Acoustic Sensors

Ultrasonic, Microphones,
Geophones, Vibrometers

Environmental Sensors

Temperature, Humidity,
Pressure, IR, etc.

Touchscreen Sensors

Capacitive, IR

Image Sensors

Thermal, Image

Biometric Sensors

Fingerprint, Heart rate, etc.

Force Sensors

Pressure, Strain

Rotation Sensors

Encoders



Sensors Metrics

Acoustic Sensors
Ultrasonic, Microphones,
Geophones, Vibrometers

Image Sensors
Thermal, Image

Motion Sensors
Gyroscope, Radar,
Accelerometer

Models



End-to-end **TinyML** application design


Datasets Preprocessing

Sound

Vision

Vibration

Quantization Pruning



Resource constraints



End-to-end **TinyML** application design


Sound



Vibration



Vision




TinyML Application

Example


Sound



Vibration




Vision



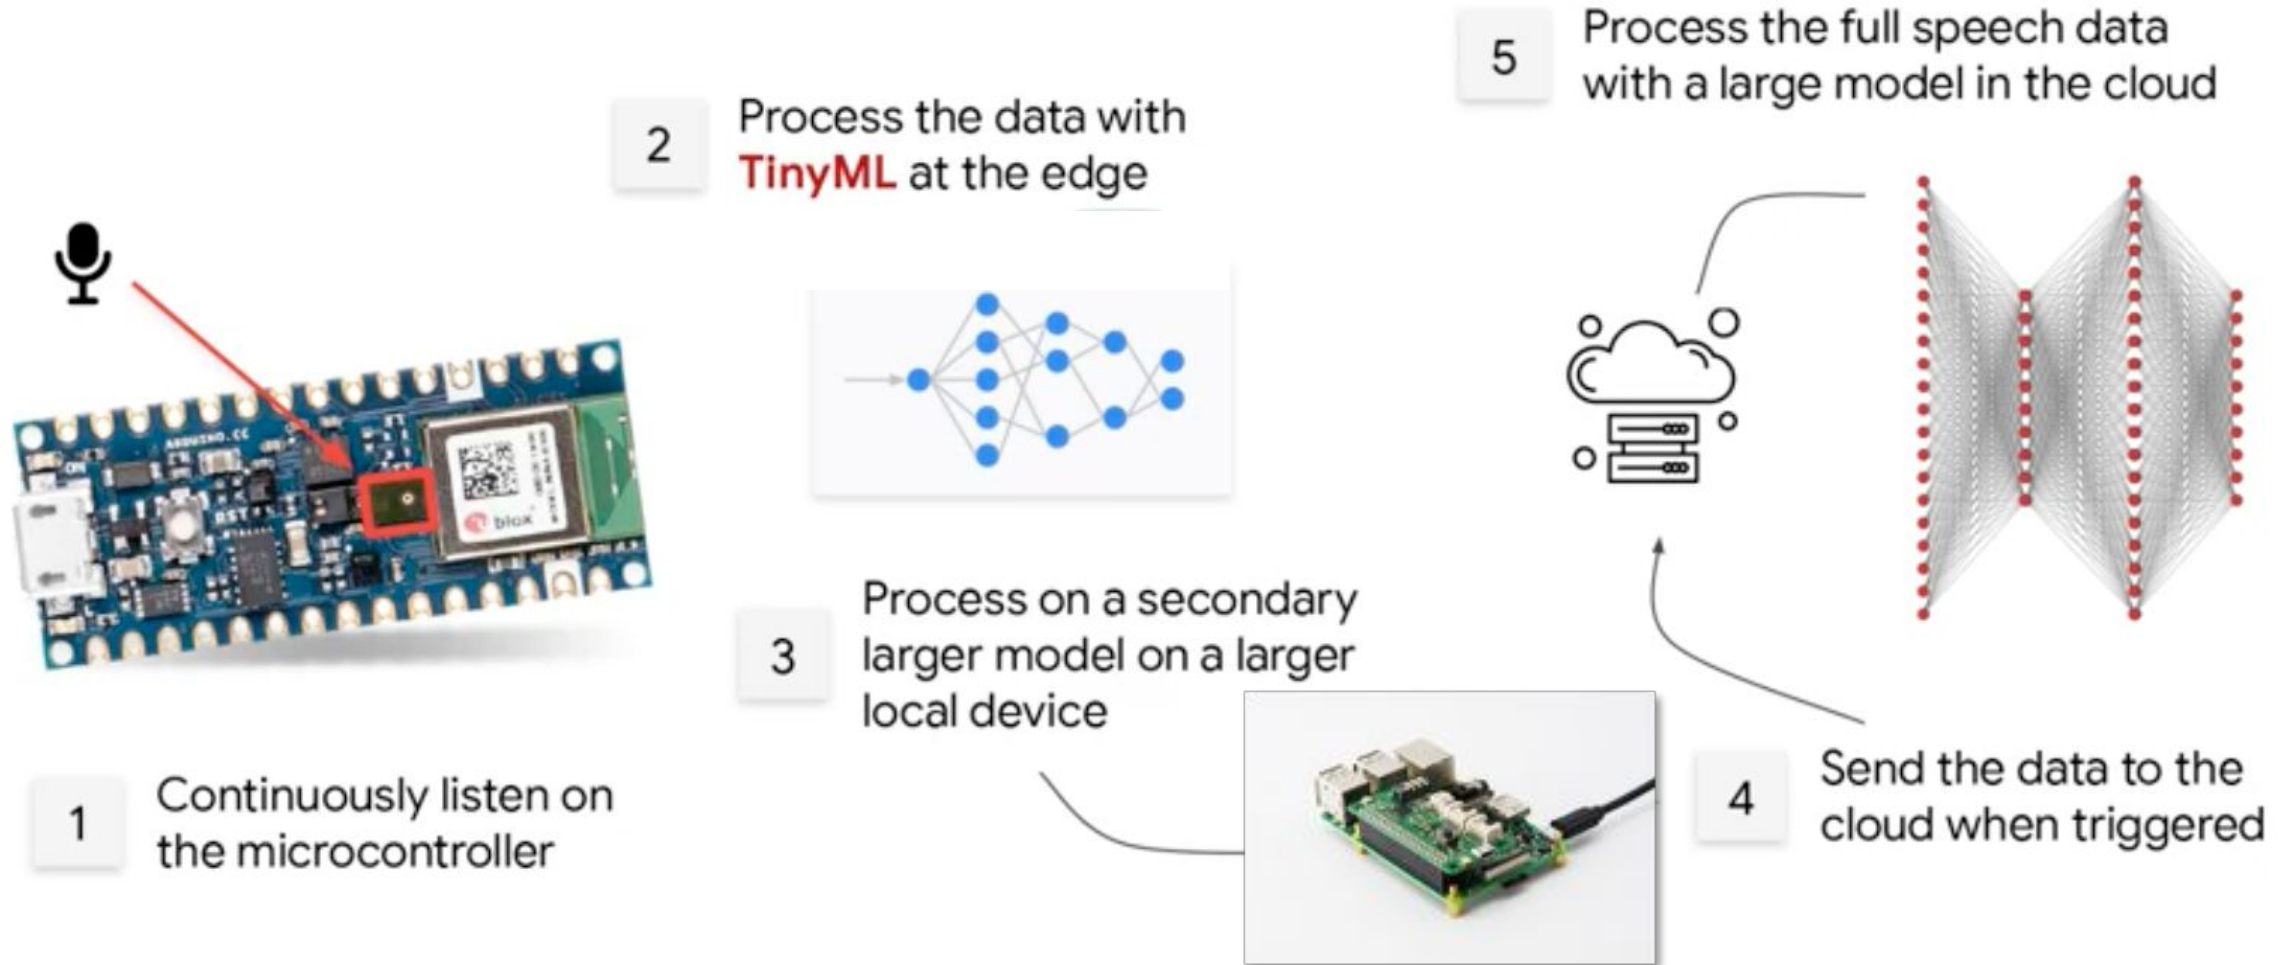
Personal Assistant





Personal Assistant



“Cascade” Detection: multi-stage model




KeyWord Spotting (KWS)



<https://mrobot.org/2021/01/27/building-an-intelligent-voice-assistant-from-scratch/>



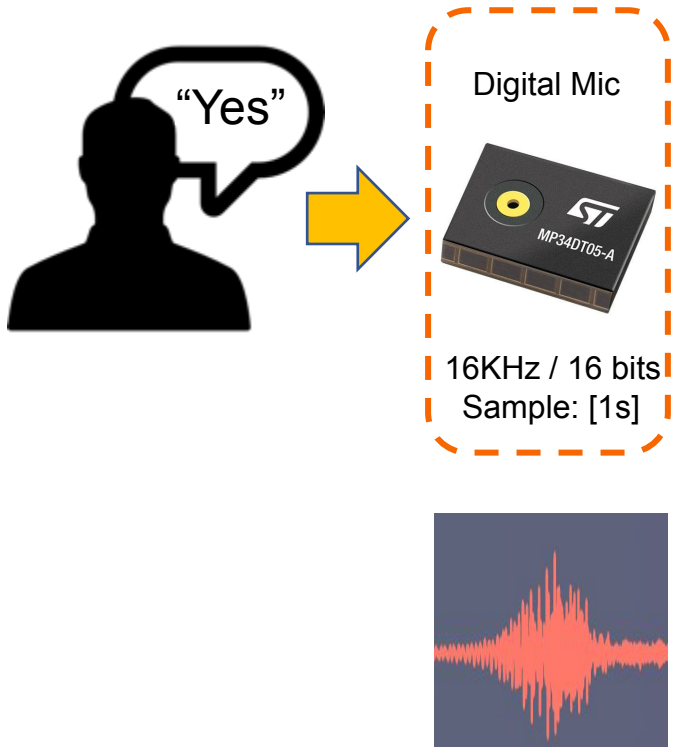
Sound



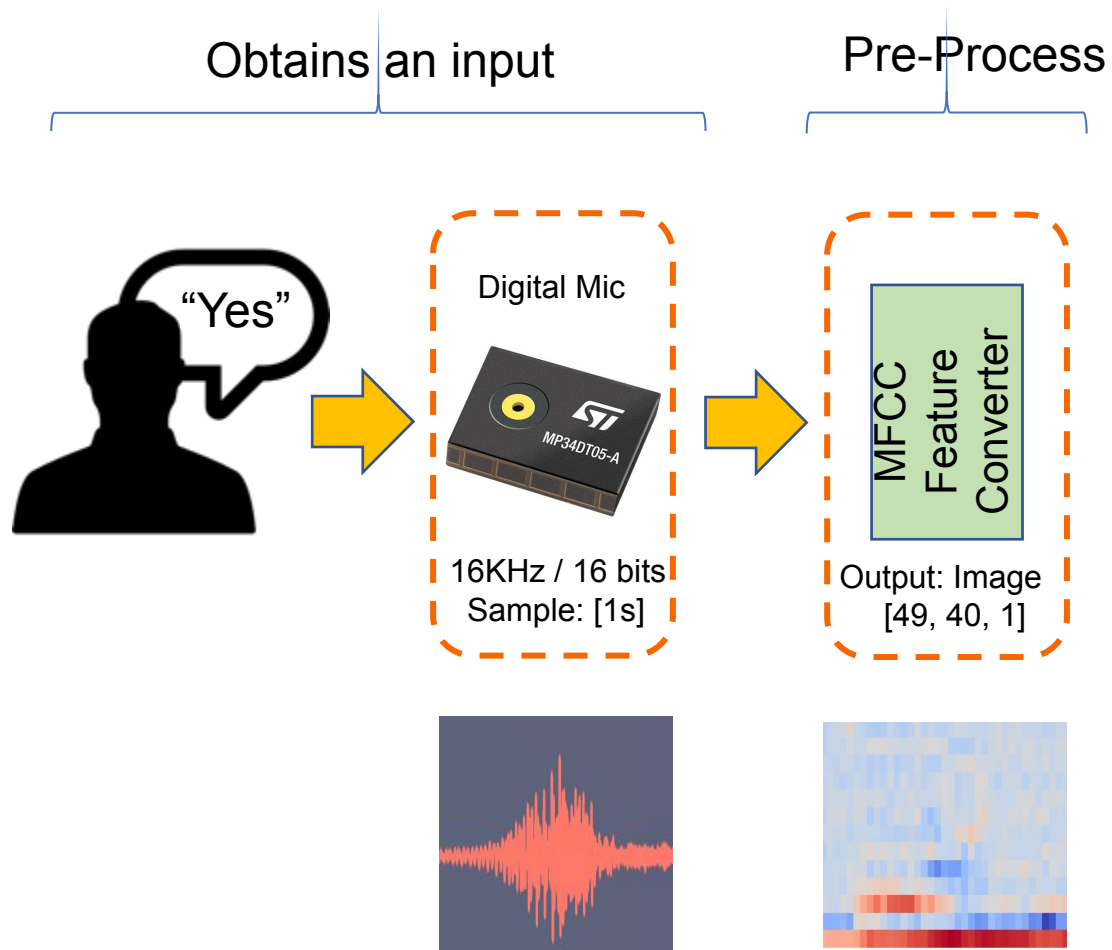
Image

KeyWord Spotting (KWS) - Inference

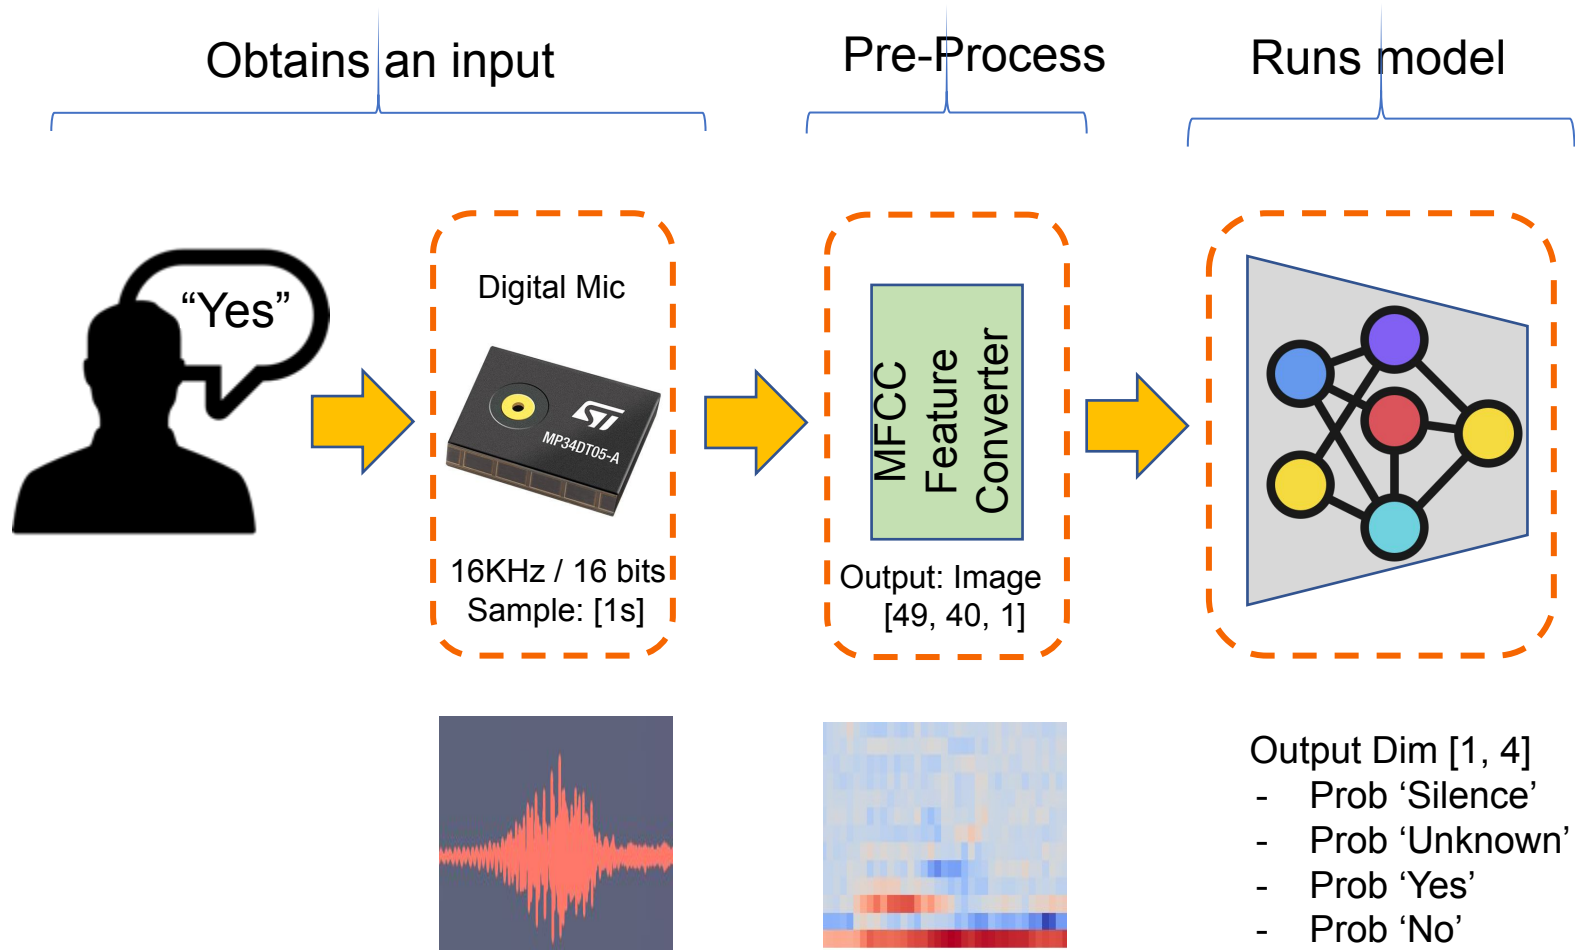
Obtains an input



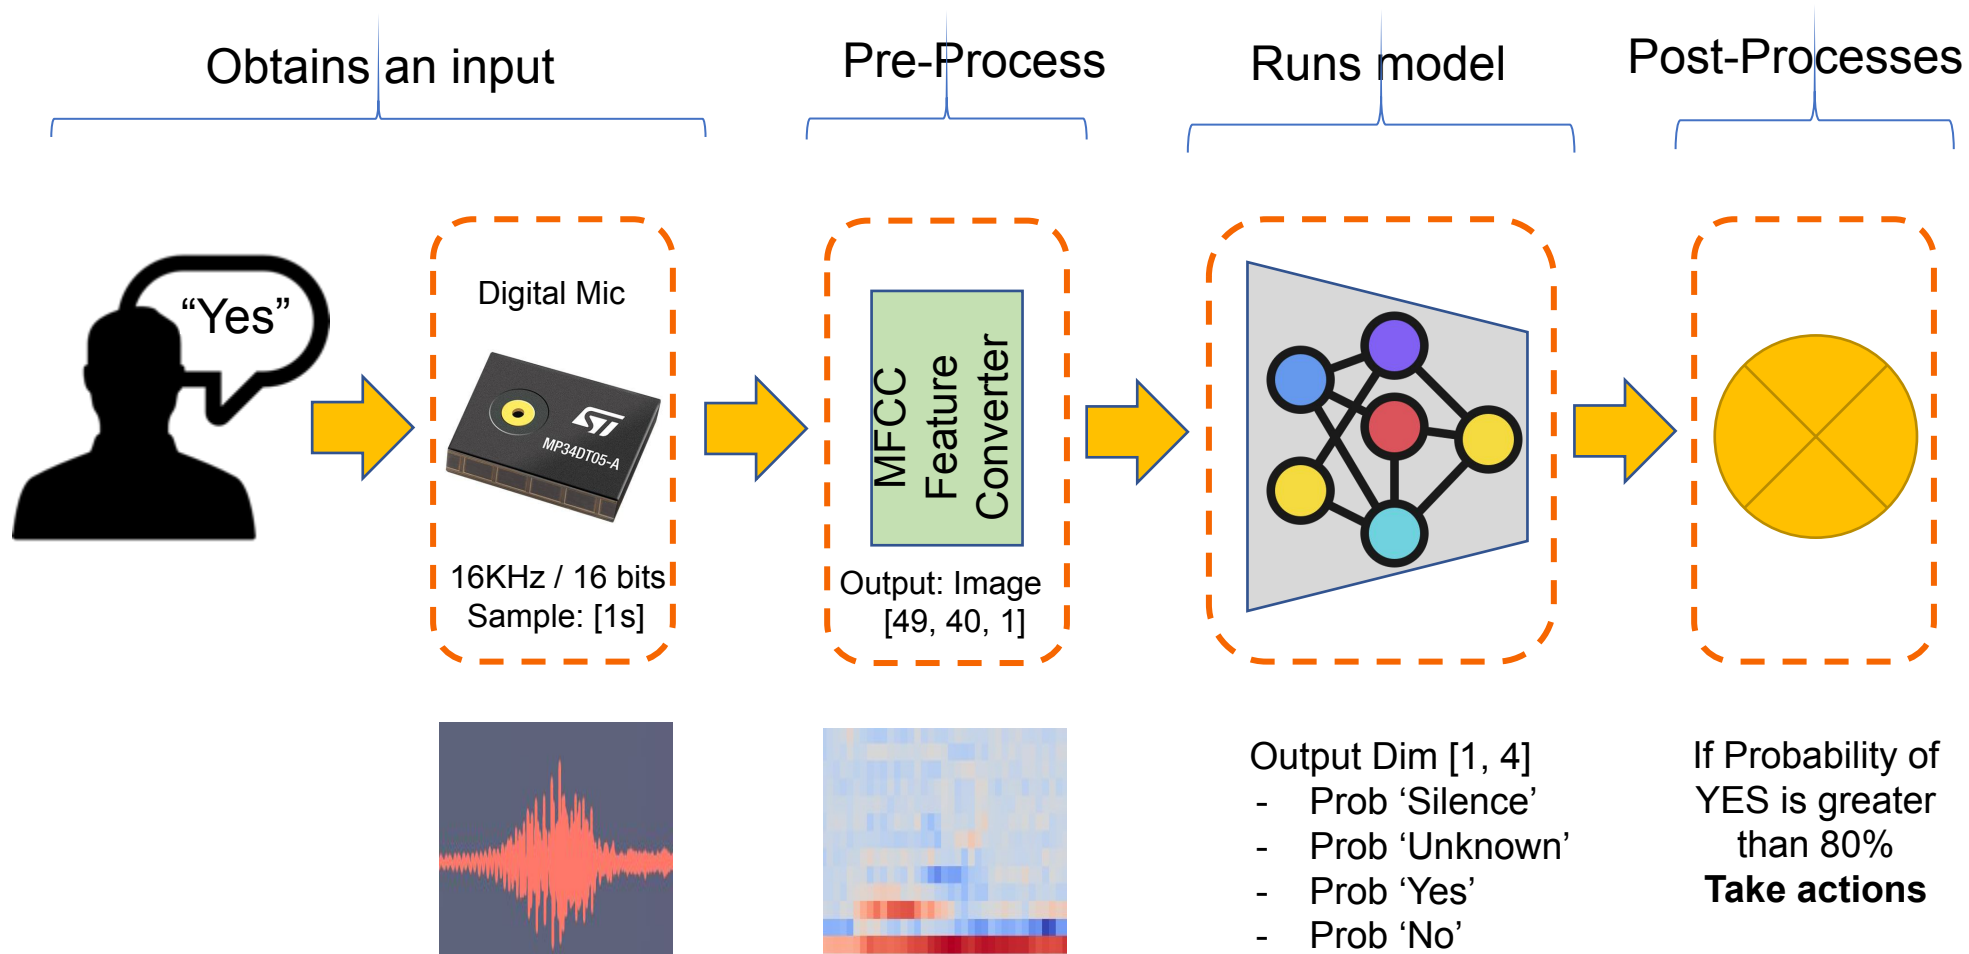
KeyWord Spotting (KWS) - Inference



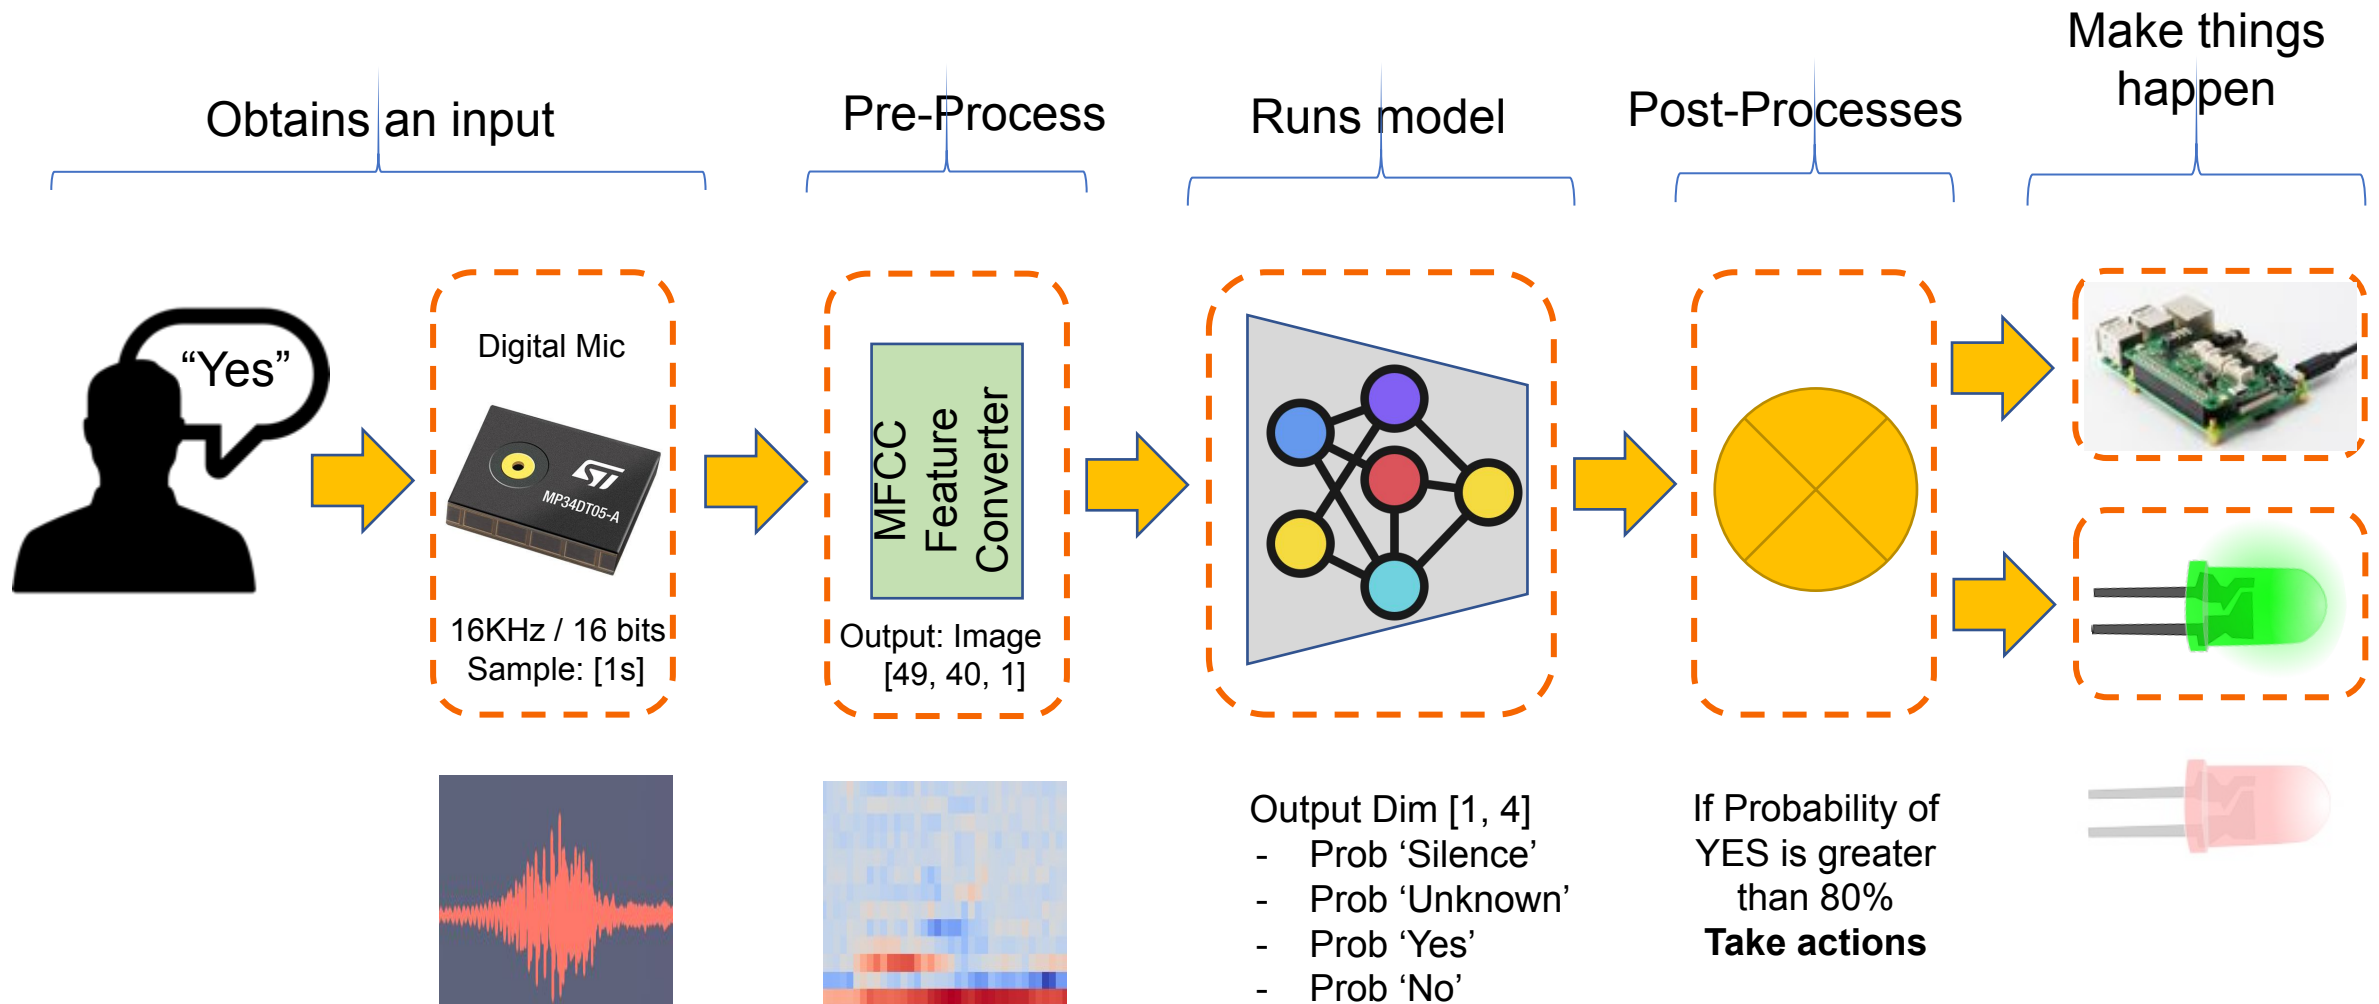
KeyWord Spotting (KWS) - Inference




KeyWord Spotting (KWS) - Inference



KeyWord Spotting (KWS) - Inference



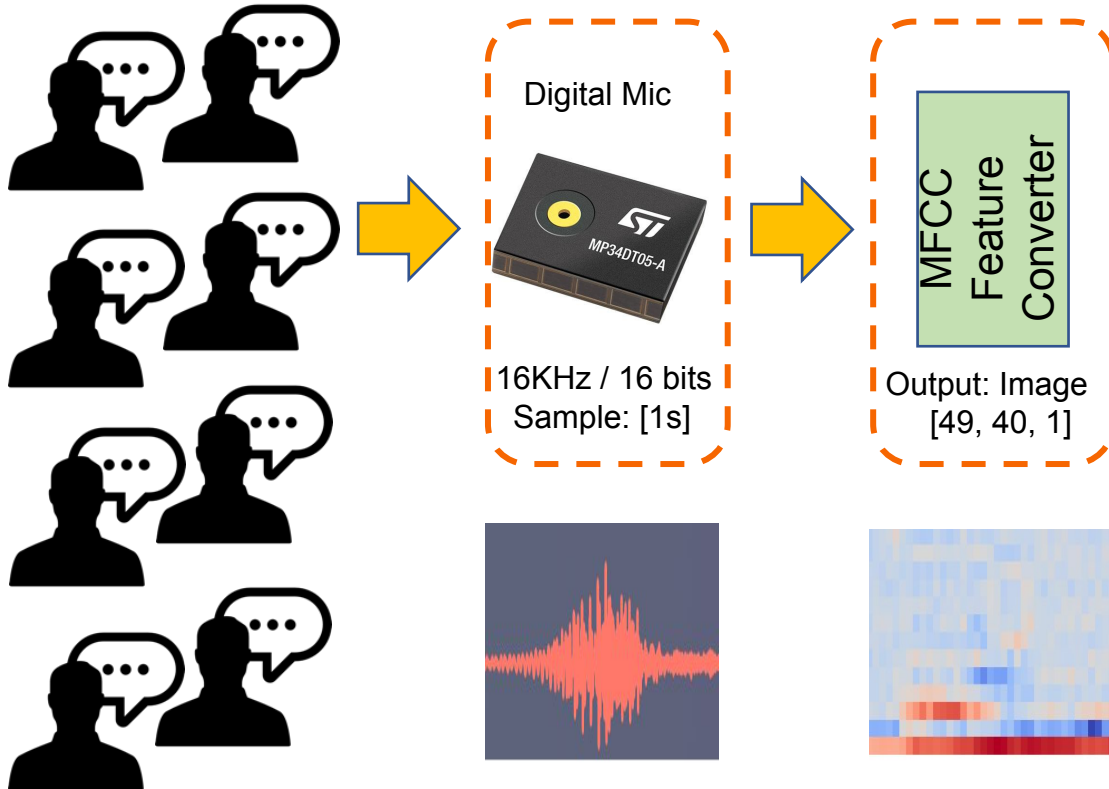
KeyWord Spotting (KWS) - Model



KeyWord Spotting (KWS) – Create Model (Training)

Obtains data

Pre-Process



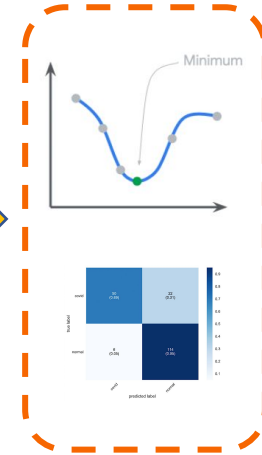






KeyWord Spotting (KWS) – Create Model (Training)

Obtains data

Pre-Process

Train model

Evaluate Model



KeyWord Spotting (KWS) – Create Model (Training)

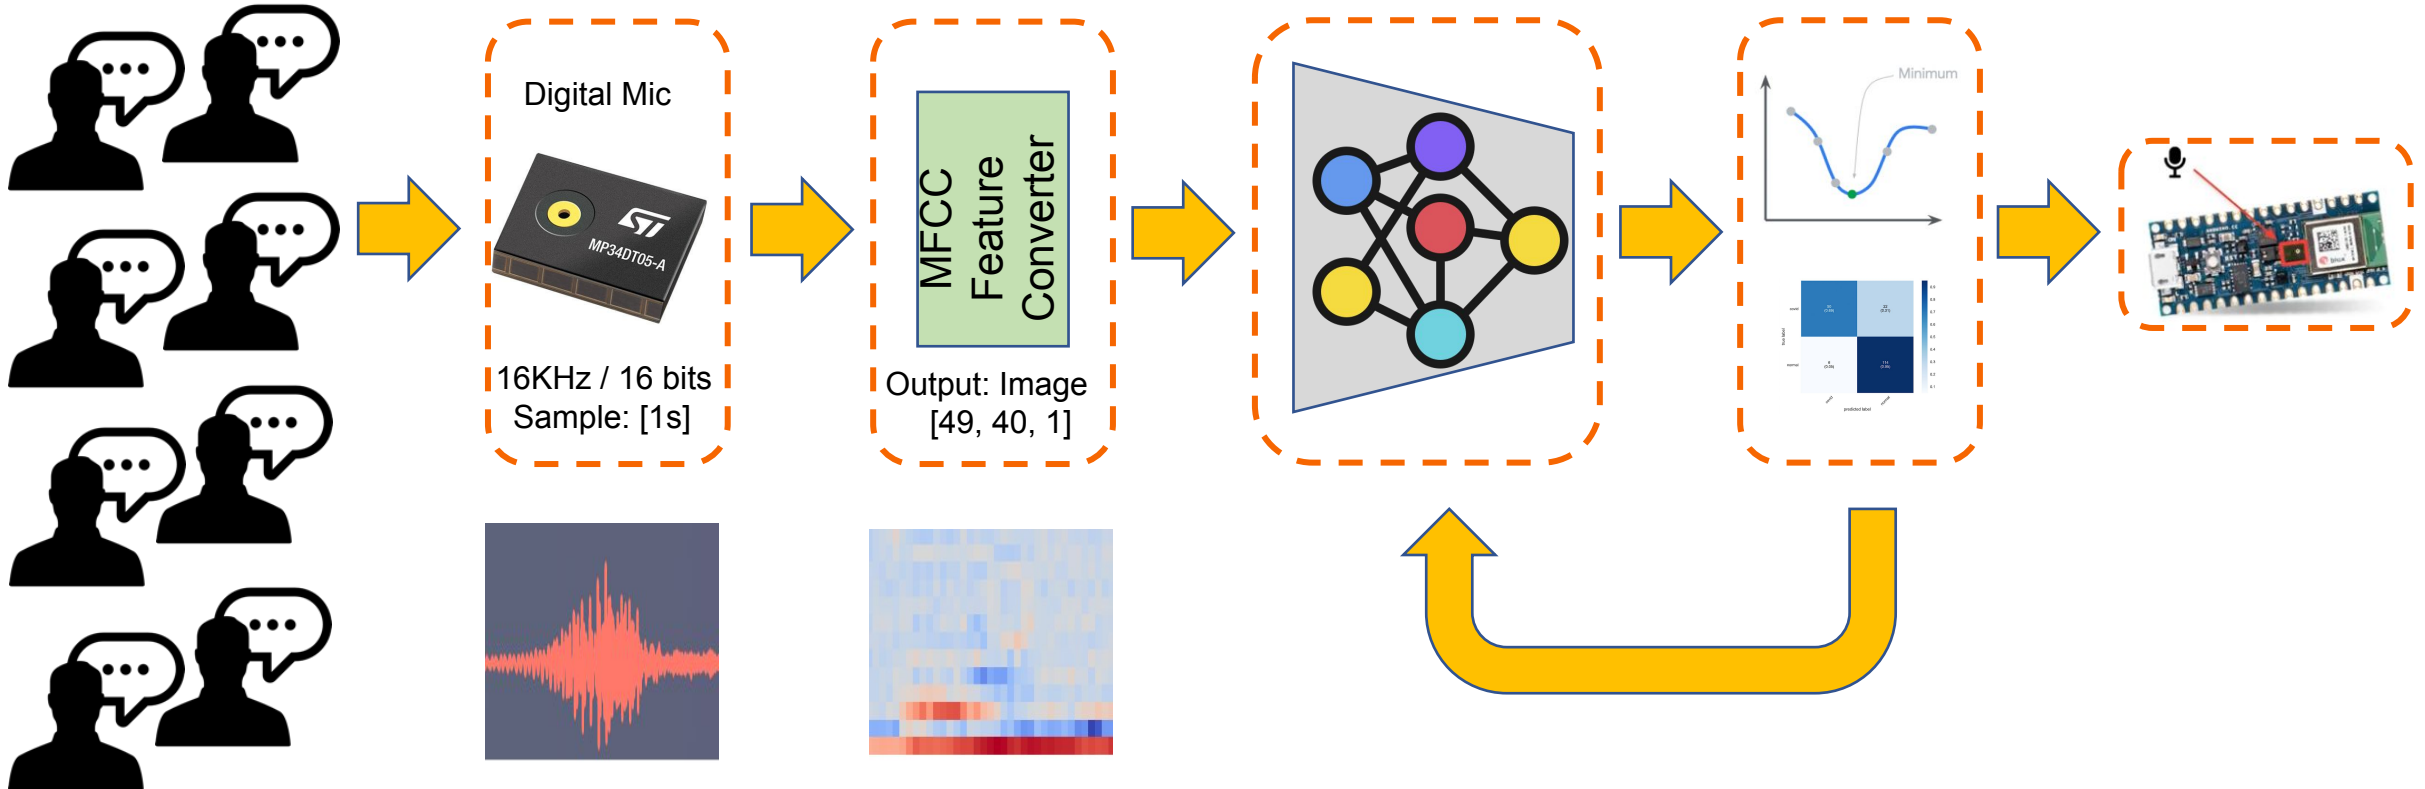
Obtains data

Pre-Process

Train model

Evaluate Model

Deploy



Reading Material

Main references

- [Harvard School of Engineering and Applied Sciences - CS249r: Tiny Machine Learning](#)
- [Professional Certificate in Tiny Machine Learning \(TinyML\) – edX/Harvard](#)
- [Introduction to Embedded Machine Learning \(Coursera\)](#)
- [Text Book: "TinyML" by Pete Warden, Daniel Situnayake](#)

I want to thank Shawn Hymel and Edge Impulse, Pete Warden and Laurence Moroney from Google, and especially Harvard professor Vijay Janapa Reddi, Ph.D. student Brian Plancher and their staff for preparing the excellent material on TinyML that is the basis of this course at UNIFEI.

The IESTI01 course is part of the TinyML4D, an initiative to make TinyML education available to everyone globally.

Thanks
And stay safe!



UNIFEI